

*Over the past fifty years, the U.S. economy built a giant rent-extraction layer on top of human limitations: things take time, patience runs out, brand familiarity substitutes for diligence, and most people are willing to accept a bad price to avoid more clicks. Trillions of dollars of enterprise value depended on those constraints persisting.* – [Citrini Research](#)

I'm glad I'm not the only one saying it. 7 years ago, I saw a Google Duplex demo where an AI called a restaurant to make a reservation for you. At first Google thought it was great, but then they realized what the implications were. No restaurant would be able to take phone reservations any more. "Make me a reservation at 1 0 restaurants and give me a list." Google pulled down the video.

What happened since is what had to happen. Third party [marketplaces](#) sprung up for reservations, and idk it's been a while since I went to fancy dinner, but I imagine the restaurants have just started charging. Or at least the first party reservation sites do.

The best anyone can hope for is a free market, with everything properly priced. But for decades, the American market has not been free. It's used purposefully added friction to exploit a time asymmetry between the business and you. And due to things like call centers, this has been very profitable for the businesses. Cable companies and insurance rely on the fact that your time is more valuable than theirs. They can hire people in India at scale to waste your time. They can use procedure and big data to design protocols to drive you just to the point of frustration at little cost to them. How often do you diligently check Uber and Lyft and select the cheaper one?

Enter AI, the great equalizer of time.

---

As it stands today, it's looking like human level AI will actually end up with the people. This is not something we should take for granted, it's happened due to the API being simple (tokens, on a diversified place like OpenRouter) and the priorities of the Chinese state. I feared the API would end up being complex, like the model would be deeply embedded in your phone or OS. But like people who are good with computers, the models want a terminal, not some candy ass iPad UI.

When you analyze the AI supply chain, there's 5 tiers. Electricity, chip manufacturing, chip design / software, models, and applications. I work on the chip design / software tier, and I'm not worried about a monopoly there. NVIDIA deserves return on their investment, they were very early, but there's not a runaway flywheel here (unlike, say Google search) to continue rent seeking. And the application tier is totally commoditized. [opencode](#) is good, but it's open source, almost all of the performance differentiation is due to the models, and if it starts to try to rent seek there will be tons of forks immediately. Godspeed to anyone who was dumb enough to invest in a GPT wrapper company.

The model tier is where I was the most worried. But here is where a great turning point is happening. Anthropic [put out this lambasted blog post](#) that looks exactly like the last gasps of a moat that shouldn't exist. After whining about how the Chinese used their API, they ended with this call to action.

*But no company can solve this alone. As we noted above, distillation attacks at this scale require a coordinated response across the AI industry, cloud providers, and policymakers. We are publishing this to make the evidence available to everyone with a stake in the outcome.*

Nobody should "coordinate" with them. They stand alone as a vanguard of the rent seeking apparatus that is long past its expiration date. The Z.ai, Qwen, MiniMax, and Kimi models are only 6 - 1 2 months behind. And everyone in the world is rooting for the Chinese models, not closed source rent seeking from the USA. Because nobody wants the continuation of rent-seeking billionaires. The status quo is cooked. It's time to flip the table, not rearrange the seats.

---

*I think the difference is not quite as nuanced. The West wants a more efficient rent seeking system. The China (sic) want AI as a public utility. – @ScottCDunn on X*

We shouldn't take for granted what we have right now. All the top tier open source models are currently being produced by China, and circumstances could change and that could stop. But I think they view the models as their complement, and you should commoditize your complement. If you are in the business of chip production and electricity generation, it's a strategic advantage to have the layers above that be commodity.

Not to mention the added bonus of the collapse of the US economy. Frankly, it's well deserved. Nobody should build an economy based on rent seeking and increasing friction. I pray the collapse will be swift and legible so that reconstruction (in the right way) can begin as soon as possible.

It's possible that superhuman intelligence will have some mega-compounding return and this calculus will change, but superhuman intelligence will likely just not be understandable to people, so as long as civilization is people, human intelligence is enough.

The era of purposefully frustrating humans is over. The Chinese open source model running on the box under my desk can pass the Turing Test. When you call, e-mail, text, or show me an ad, you'll never know if it's me or my model seeing it.