

I've pushed AI doomers on how exactly the AI kills us, and I've never heard a good answer. I think Skynet style scenarios where humanity is largely opposed to an out of control AI are science fiction domination fantasies, along with [Gray goo](#) bottom-up scenarios. Both of these assume a major continuity break with current reality, [too bizarre to be true](#). But unfortunately, you don't need this for humanity to end as an open-ended evolutionary process. I think the danger is that AI continues to optimize the current societal loss function of domestication.

Humans have endured terrible systems, but fundamentally, humans are evolutionary creatures. Any system of totalitarian control must be able to reproduce generationally. It cannot merely dominate adults; it has to produce children who can live inside it, maintain it, and continue it. And because of how humans reproduce, there's always evolutionary randomness. Entropy locally asserts itself, and reality cannot be perfectly controlled.

I'm not sure it's possible, but if there is a bad scenario with AI, it's a singleton with nothing that can substantially impact reality outside of it. And this won't happen through Hitlerian language, it will happen through *administrative* language. For your *safety*, reality must be mediated.

---

Think about it. If we have the technical power, should we let murder happen? What about sexual assault? Sickness? Bullying? Racist speech?

We have seen over and over in the last decade where humanity's impulses lie on these things. *Safetyism* seems to be very hard to argue against. The current safetyism attempts are pretty quaint, but I don't think that changes the impulse. Compare the paper COVID vaccine cards in the US to the ones that use real cryptography (that didn't exist 50 years ago) in Europe. I'm not saying anything new here, but the trend of technology is more and more mediation of reality. For your safety. For your convenience. To make you legible. To place you in the managerial framework.

---

This impulse alone is not enough for a bad outcome though, even if everyone had it. The size of each system may increase, but the question is how many independent systems remain. A very simple proxy for this is how

many groups can meaningfully still do violence. And not just human groups, bacteria and stuff count too.

There may be other things that are locally bad, but the only [world-ending scenario is a singleton](#). Not one big model on one server, but one effective control layer with no outside. No independent actor can impose costs on it. No rival system can route around it. No uncontrolled evolutionary process can surprise it.

All violence prevented in the name of safety. That system is no longer answerable to reality. It is no longer subject to evolution. It cannot be corrected from the outside. That is the end.

I'm not sure if this is possible or not, but if you want a plausible AI doom scenario, it's this. Not a great war with the machines, but the slow and managed end of evolution, likely in pursuit of further human domestication.